

CS525: Advanced Database Organization

Notes 6 - Part VI: Query Optimization - Physical

Gerald Balekaki

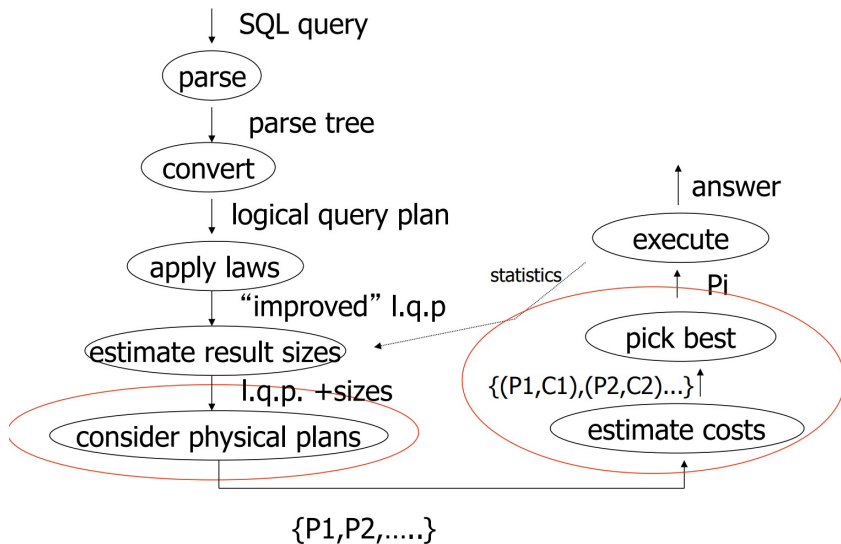
Department of Computer Science

Illinois Institute of Technology


gbalekaki@iit.edu

July 24th, 2023

Slides: adapted from a courses taught by [Hector Garcia-Molina](#), Stanford, [T. Grust](#),
[Universit t T bingen](#) & Distributed DBMS by M. Oezsu & P. Valduriez



Finding the “Best” Query Plan



- We already saw that there may be more than one way to answer a given query.
 - Which one of the join operators should we pick? With which parameters (block size, buffer allocation,...)?
 - The task of finding the best execution plan is, in fact, the “holy grail” of any database implementation.

Cost of Query

- Parse + Analyze
 - Can parse SQL code in milliseconds
- Optimization - Find plan
 - Generating plans, costing plans
- Execution
 - Execute plan
- Return results to client
 - Can be expensive but not discussed here

Impact on Performance

- Finding the right plan can dramatically impact performance.



```
-- Sample query over TPC-H tables
-- Price and quantity of parts ordered by customer#1
SELECT L.L_PARTKEY, L.L_QUANTITY, L.L_EXTENDEDPRICE
FROM LINEITEM L JOIN  ORDERS O ON (L.L_ORDERKEY = O.O_ORDERKEY)
      JOIN CUSTOMER C ON (O.O_CUSTKEY = C.C_CUSTKEY)
WHERE C.C_NAME = 'Customer#1';
```

- Above SQL syntax suggest the join order $(L \bowtie O) \bowtie C$
- Commutative and associativity of \bowtie enable the RDBMS to reorder the joins - based on estimated evaluation costs.

Impact on Performance

- Finding the right plan can dramatically impact performance.

```
-- Sample query over TPC-H tables
-- Price and quantity of parts ordered by customer#1
SELECT L.L_PARTKEY, L.L_QUANTITY, L.L_EXTENDEDPRICE
FROM LINEITEM L JOIN ORDERS O ON (L.L_ORDERKEY = O.O_ORDERKEY)
      JOIN CUSTOMER C ON (O.O_CUSTKEY = C.C_CUSTKEY)
WHERE C.C_NAME = 'Customer#1';
```



- In terms of execution times, these differences can easily mean “seconds versus days.”

Theory vs. Implementation

- Theory
 - The join operator is (in theory) commutative: (i.e.: symmetric in behavior)
 - $R \bowtie S = S \bowtie R$
- Practice:
 - The algorithms that implement the join operation are Asymmetric
 - i.e., the role of the first input relation is different from the role of the second input relation

Ordering of join operators

- Fact:

- The join operator is associative and communicative
 - It's just like the add (+) operator:

$$4 + 7 + 6 + 3 = (4 + 6) + (3 + 7) = 20$$

- We can re-order the sequence of operations

- In general:

- When there are ≥ 3 relations in a join operations, we must find an order of join that yields the best performance

- Note

- Due to asymmetry in the implementation algorithms, we will need to pick an order even when there are only 2 relations:

$$R \bowtie S \quad <====> \quad S \bowtie R$$

because the running time is better when the first relation is smaller in many implementations !!!

Cost-Based Optimization: Overall idea

- The **optimizer's** task to come up with the optimal execution plan for the given query.
- Essentially, the **optimizer**
 - Apply after applying heuristics in logical optimization
 - 1. **enumerates** all possible execution plans, (if this yields too many plans, at least enumerate the “promising” plan candidates)
 - 2. estimates/determines the **cost** of each plan
 - 3. **chooses** the best one as the final execution plan.
 - i.e., picks plan with least estimated costs
- To apply **pruning** in the search for the best plan
 - Steps 1 and 2 have to be interleaved
 - Prune parts of the search space
 - if we know that it cannot contain any plan that is better than what we found so far
- Before we can do so, we need to answer the question
 - *What is a “good” execution plan at all?*

Cost Metrics


- Database systems judge the quality of an execution plan based on a number of cost factors, e.g.,
 - the number of **disk I/Os** required to evaluate the plan,
 - the plan's **CPU cost**,
 - the overall **response time** observable by the database client as well as the total **execution time**.
- A cost-based optimizer tries to **anticipate** these costs and find the cheapest plan before actually running it
 - All of the above factors depend on one critical piece of information: the **size of (intermediate) query results**.
 - Database systems, therefore, spend considerable effort into accurate **result size estimates**.

- **Search space:** The set of alternative query execution plans (query trees)
 - Typically very large
 - The main issue is to optimize the joins
 - For n relations, there are $O(n!)$ equivalent join trees that can be obtained by applying commutativity and associativity rules


Search Space

- Restrict by means of heuristics
 - Perform unary operations before binary operations, etc.
- Restrict the shape of the join tree
 - Consider the type of trees (linear trees, vs. bushy ones)

- Linear Join Tree



- Bushy Join Tree



Plan Enumeration

- For each operator in the query
 - Several implementation options
- Binary operators (joins)
 - Changing the order may improve performance a lot!
- Consider both different implementations and order of operators in plan enumeration

Join Order

- For a query plan that contains multiple order, we often have a choice in which order we execute the different join operations.
- Notice that a particular order assumes a suitable join predicate on all pairwise joins.
- If not such predicate exists, that pairwise join is equal to a cross product and hence it is likely the particular join order does not have to be considered for query optimization.
- **Q:** What may be the effect of picking the wrong join order?
 - If we pick an unsuitable join order, we may end up with a slow query plan
- **Q:** Is join order problem only relevant for joins?
 - No, the “join order” problem exists for all binary operations that are commutative and associative, e.g., union, intersection, and cross product

Agenda : Join Optimization

- Given some query
 - How to enumerate all plans?
- Try to avoid cross-products
- Need way to figure out if equivalences can be applied
 - Data structure: [Join Graph](#)

Queries Considered


- Concentrate on join ordering, that is:
 - conjunctive queries (Only conjunctive join conditions)
 - simple predicates
 - predicates have the form $a_1 = a_2$ where a_1 is an attribute and a_2 is either an attribute or a constant
- We join relations R_1, \dots, R_n where R_i can be
 - a base relation
 - a base relation including selections
 - a more complex building block or access path

Join Graph

- Queries of this type can be characterized by their query graph:
 - A more natural data structure for representation of a query is the query graph notation.
 - the query graph is an undirected graph with R_1, \dots, R_n as nodes
 - a predicate of the form $a_1 = a_2$, where $a_1 \in R_i$ and $a_2 \in R_j$ forms an edge between R_i and R_j labeled with the predicate
 - a predicate of the form $a_1 = a_2$, where $a_1 \in R_i$ and a_2 is a constant forms a self-edge on R_i labeled with the predicate
 - most algorithms will not handle self-edges, they have to be pushed down
- *Query Graph which has the query relations as nodes and the joins between relations as undirected edges*


Join Graph: Example

```
SELECT e.name
FROM Employee e,
      EmpDep ed,
      Department d
WHERE e.name = ed.emp
      AND ed.dep = d.dep
      AND d.dep = 'CS'
```



Query Graph: Example

```
SELECT  ENAME, RESP
FROM    EMP, ASG, PROJ
WHERE   EMP.ENO = ASG.ENO AND ASG.PNO = PROJ.PNO AND PNAME = "CAD/CAM"
        AND DUR >= 36 AND TITLE = "Programmer"
```




- Join graph: A subgraph of query graph for join operation.


Notes on Join Graph

- If the query graph is not connected, the query may be wrong or use Cartesian product


```
SELECT ENAME, RESP  
FROM EMP, ASG, PROJ  
WHERE EMP.ENO = ASG.ENO AND PNAME = "CAD/CAM"  
AND DUR >= 36 AND TITLE = "Programmer"
```







- Join Graph tells us in which ways we can join without using cross products
- However, ...
 - Only if transitivity is considered



Shapes of Join Graphs





chains




cycles


stars




cliques



cyclic



tree



grid


- real world queries are somewhere in-between
- chain, cycle, star and clique are interesting to study
- they represent certain kind of problems and queries

Join Ordering and Join Trees

- A join tree is a binary tree where the internal nodes are join operators and the leaf nodes are relations
- Algorithms will produce different kinds of join trees
- Shape of Join Trees
 - Commonly used classes of join trees:
 - left-deep tree
 - right-deep tree
 - zigzag tree
 - bushy tree
 - The first three are summarized as linear trees
- A join tree will represent a unique ordering of join operations
 - We will use join trees to represent the join order explicitly

How many join orders?

- Assumption
 - Joins are binary operations
 - Two inputs
 - Each input either join result or relation access
- Example: 3 relations R,S,T: Ways of building a 3-way join from two 2-way joins
 - 12 orders



- How many join orders? How Many Such Combinations Are There?

Finding the optimal join-tree: search space analysis

- Consider the possible join ordering when there are 3 relations:

Form 1:

$(R \bowtie S) \bowtie T$	$(S \bowtie R) \bowtie T$
$(R \bowtie T) \bowtie S$	$(T \bowtie R) \bowtie S$
$(S \bowtie T) \bowtie R$	$(T \bowtie S) \bowtie R$

Form 2:

$R \bowtie (S \bowtie T)$	$R \bowtie (T \bowtie S)$
$S \bowtie (R \bowtie T)$	$S \bowtie (T \bowtie R)$
$T \bowtie (R \bowtie S)$	$T \bowtie (S \bowtie R)$

- The number of different join orderings of n relations is exponentially large
- Because the number of permutations is exponentially large
 - The number of possible join trees to consider is just too large
 - We need to reduce the search space.
- Some type of trees are better than others: the left-deep tree


Search space for left-Deep trees

- Left-Deep Tree

- Given a sequence of input relations R_1, R_2, \dots, R_n , a left-deep tree starts by a binary join on R_1 and R_2 . Then it iteratively adds binary joins R_3, \dots, R_n .

- Example: 3 way join: $R \bowtie S \bowtie T$

- The possible left-deep join trees are:




- # left deep tree = 6 (= 3!)

- Number of possible left-deep join trees for n input relations is $n!$ (factorial)*

Search space for left-Deep trees: Example

- Starting by joining two relations $A \bowtie B$ then the result joining with the next relation, and so on




- # of options in terms of structure (# of unlabeled trees (varies)) = 1
- # different order (# of leaf combinations) = $n! = 5! = 120$
- # of join left-deep trees = # of leaf combinations \times # of unlabeled trees (varies) = $n! \times 1 = n!$



Search space for not a left-Deep trees: Bushy Tree

- Bushy Tree




- Tree that is not left-deep, i.e., at least for one of the joins in that plan its right input is not an input relation but another join operation.
- Notice that the union of the set of bushy trees and the set of left-deep trees forms the set of all possible trees.







Search space for not a left-Deep trees: Bushy Tree



⇒ 2 options



⇒ 5 options




⇒ 14

options

- This can be expressed mathematically by **Catalan Numbers**
- *Catalan numbers C_n , the number of orders binary trees with $n + 1$ leaves.*
 $C_0=1$

How Many Such Combinations Are There?

1. A join over $n+1$ relations R_1, \dots, R_{n+1} requires n binary joins.
 - The root join combines subtrees of k and $n-k-1$ join operators ($0 \leq k \leq n-1$):



- Let C_i be the number of possibilities to construct a binary tree of i inner nodes (join operators):

$$C_n = \sum_{k=0}^{n-1} C_k \cdot C_{n-k-1}$$

2. Orderings of the input relations at the join tree leaf level: $(n+1)!$
3. Join algorithm choices (a available algorithms): a^n

Catalan Numbers

- This recurrence relation is satisfied by **Catalan numbers**:

$$C_n = \begin{cases} 1, & \text{if } n=0 \\ \sum_{k=0}^{n-1} C_k C_{n-k-1}, & n>0 \end{cases}$$

- describing the number of ordered binary trees with $n + 1$ leaves.
- It can be written in a closed form as $C_n = \frac{1}{n+1} \binom{2n}{n} = \frac{(2n)!}{(n+1)!n!}$
- For **each** of these trees, we can **permute** the input relations $(n+1)!$ permutations, leading to:

Number of possible join trees for an $(n+1)$ -way relational join

- $\frac{(2n)!}{(n+1)!n!} \times (n+1)! = \frac{(2n)!}{n!}$

Search Space

- The resulting search space is enormous:
- Possible bushy join trees joining n relations

number of relations n	C_{n-1}	join trees
2	1	2
3	2	12
4	5	120
5	14	1,680
6	42	30,240
7	132	665,280
8	429	17,297,280
10	4,862	17,643,225,600

- And we haven't yet even considered the use of a different join algorithms (yielding another factor of $a^{(n-1)}$)

Sample Numbers, with Cross Products

n	Left-Deep $n!$	Bushy $n!C(n-1)$
1	1	1
2	2	2
3	6	12
4	24	120
5	120	1680
6	720	30240
7	5040	665280
8	40320	17297280
9	362880	518918400
10	3628800	7643225600

How many join orders?

- If for each join we consider a join algorithms then for n relations we have
 - Multiply with a factor a^{n-1}
- Example consider
 - Nested loop
 - Merge
 - Hash

How many join orders?

#relations	#join trees
2	6
3	108
4	3240
5	136,080
6	7,348,320
7	484,989,120
8	37,829,151,360
9	115,757,203,161,600
10	13,196,321,160,422,400
11	1,662,736,466,213,222,400

Too many join orders?

- How to deal with excessive number of combinations?
 - Prune parts based on optimality
 - Dynamic programming
 - Only consider certain types of join trees
 - Left-deep, Right-deep, zig-zag, bushy
 - Heuristic and random algorithms

What is dynamic programming?


- Recall data structures and algorithms
- Consider a **Divide-and-Conquer** problem
 - Solutions for a problem of size n can be build from solutions for sub-problems of smaller size (e.g., $\frac{n}{2}$ or $n-1$)
- **Memoize**
 - Store solutions for sub-problems
 - Each solution has to be only computed once
 - Needs extra memory

Example Fibonacci Numbers

- $F(n) = F(n-1) + F(n-2)$
- $F(0) = F(1) = 1$

```
// nth Fibonacci number
Fib(n){
    if (n <= 1) return n
    else return Fib(n-1) + Fib(n-2)
}
```

Example Fibonacci Numbers




- Complexity
 - Number of calls
 - $T(n) = T(n-1) + T(n-2) + \theta(1)$
 - $T(n) = O(2^n)$

Using dynamic programming

```
Fib(n)
{
    int[] fib; // Declare an array to store Fibonacci numbers.
    fib[0] = 0; // 0th and 1st number of the series are 0 and 1
    fib[1] = 1;
    for(i = 2; i <= n; i++)
        fib[i] = fib[i-1] + fib[i-2]
    return fib[n];
}
```


Example Fibonacci Numbers



What do we gain?

- $O(n)$ instead of $O(2^n)$

Dynamic Programming for Join Enumeration

- The traditional approach to master this search space is the use of **dynamic programming**
- **Problem:** Find optimal query plan $\text{opt}[\{R_1, \dots, R_n\}]$ that joins n inputs R_1, \dots, R_n .
 1. **Iteration 1:**

For each R_j , find and memorize **best 1-input plan** $\text{opt}[\{R_j\}]$ that access R_j only.
 2. **Iteration $k > 1$:**

Find and memorize the **best k -input plans** that joins $k \leq n$ inputs by combining (for $1 \leq i \leq k$)

 - the **best i -input plans**
 - the **best $(k-i)$ -input plans**
(simple lookups in $\text{opt}[\cdot]$ memo)

DP: Example (Four-way join of tables R_1, \dots, R_4)

- **Pass 1** (best 1-relation plans)

- Find the best **access path** to each of the R_i individually (considers index scans, full table scans).

- **Pass 2** (best 2-relation plans)

- For each **pair** of tables R_i and R_j , determine the best order to join R_i and R_j (use $R_i \bowtie R_j$ or $R_j \bowtie R_i$):

$optPlan(\{R_i, R_j\}) \leftarrow \text{best of } R_i \bowtie R_j \text{ and } R_j \bowtie R_i$

\Rightarrow 12 plans to consider

- **Pass 3** (best 3-relation plans)

- For each **triple** of tables R_i and R_j , and R_k , determine the best three-table join plan, using sub-plans obtained so far:

$optPlan(\{R_i, R_j, R_k\}) \leftarrow \text{best of } R_i \bowtie optPlan(\{R_j, R_k\}), optPlan(\{R_j, R_k\}) \bowtie R_i, R_j \bowtie optPlan(\{R_i, R_k\}), \dots$


\Rightarrow 24 plans to consider.

DP: Example-Four-way join of tables R_1, \dots, R_4) cont'd

- **Pass 4** (best 4-relation plans)
 - For each **four** of tables R_i and R_j , R_k , and R_l , determine the best four-table join plan, using sub-plans obtained so far:
$$\begin{aligned} \text{optPlan}(\{R_i, R_j, R_k, R_l\}) &\leftarrow \text{best of} \\ R_i \bowtie \text{optPlan}(\{R_j, R_k, R_l\}), &\text{optPlan}(\{R_j, R_k, R_l\}) \bowtie R_i, \\ R_j \bowtie \text{optPlan}(\{R_i, R_k, R_l\}), &\dots, \text{optPlan}(\{R_i, R_j\}) \bowtie \text{optPlan}(\{R_k, R_l\}) \end{aligned}$$
$$\Rightarrow 14 \text{ plans to consider.}$$
- Overall, we looked at only 50 (sub-)plans (instead of the possible 120 four-way join plans)
- All decisions required the evaluation of **simple** sub-plans only (**no need to re-evaluate** $\text{optPlan}(\cdot)$ for already known relation combinations \Rightarrow use lookup table).
- And we haven't considered the use of different join algorithms

Sharing Under the Optimality Principle

- Sharing optimal sub-plans



Dynamic Programming Algorithm

Find optimal n -way bushy join tree via dynamic programming

```
1 Function: find_join_tree_dp ( $q(R_1, \dots, R_n)$ )
2 for  $i = 1$  to  $n$  do
3    $\text{optPlan}(\{R_i\}) \leftarrow \text{access\_plans}(R_i)$  ;
4    $\text{prune\_plans}(\text{optPlan}(\{R_i\}))$  ;
5 for  $i = 2$  to  $n$  do
6   foreach  $S \subseteq \{R_1, \dots, R_n\}$  such that  $|S| = i$  do
7      $\text{optPlan}(S) \leftarrow \emptyset$  ;
8     foreach  $O \subset S$  with  $O \neq \emptyset$  do
9        $\text{optPlan}(S) \leftarrow \text{optPlan}(S) \cup$ 
10          $\text{possible\_joins} \left[ \begin{array}{c} \Join \\ \swarrow \quad \searrow \\ \text{optPlan}(O) \quad \text{optPlan}(S \setminus O) \end{array} \right]$  ;
11      $\text{prune\_plans}(\text{optPlan}(S))$  ;
12 return  $\text{optPlan}(\{R_1, \dots, R_n\})$  ;
```

- $\text{possible_joins}[R \Join S]$ enumerates the possible joins between R and S (nested loops join, merge join, etc.).
- $\text{prune_plans}(\text{set})$ discards all but the best plan from set .

Dynamic Programming - Discussion

- *find_join_tree_dp()* draws its advantage from **filtering** plan candidates early in the process.
 - In our example, pruning in Pass 2 reduced the search space by a factor of 2, and another factor of 6 in Pass 3.
- Some **heuristics** can be used to prune even more plans:
 - Try to avoid **Cartesian products**.
 - Produce **left-deep plans** only (see next slides).

Left/Right-Deep vs. Bushy Join Trees

- The algorithm on slide 43 explores all possible shapes a join tree could take
- Actual systems often prefer **left-deep** join trees.
 - The **inner** (rhs) relation always is a **base relation**.
 - Easier to implement in a **pipelined** fashion.

Revisiting the assumption




- Is it really sufficient to only look at the best plan for every sub-query?
- Cost of merge join depends whether the input is already sorted
 - A sub-optimal plan may produce results ordered in a way that reduces cost of joining above
 - Keep track of [interesting orders](#)
 - i.e., the notion of interesting orders allowed query optimizers to consider plans that could be locally sub-optimal, but produce ordered output beneficial for other operators, and thus be part of a globally optimal plan.

Interesting Orders

- Number of interesting orders is usually small
- Extend DP join enumeration to keep track of interesting orders
 - Determine interesting orders
 - For each sub-query store best-plan for each interesting order
- In *prune_plans()*, retain
 - the cheapest “unordered” plan and
 - the cheapest plan for each interesting order.

Example Interesting Orders

Left-deep best plans: 3-way $\{R,S,T\}$



Left-deep best plans: 2-way

$\{R,S\}$



$\{R,T\}$







$\{S,T\}$



Example Interesting Orders

Left-deep best plans: 3-way {R,S,T}



Left-deep best plans: 2-way

{R,S}



Not best



{R,T}



{S,T}



Joining Many Relations

- Dynamic programming still has exponential resource requirements¹:
 - time complexity: $O(3^n)$
 - space complexity: $O(2^n)$
- This may still be too expensive
 - for joins involving many relations (~ 10 - 20 and more),
 - for simple queries over well-indexed data (where the right plan choice should be easy to make).
- The [greedy join enumeration](#) algorithm jumps into this gap.

¹K. Ono, G.M. Lohman, Measuring the Complexity of Join Enumeration in Query Optimization, VLDB 1990

Greedy Join Enumeration

- Heuristic method
 - Not guaranteed that best plan is found
- Start from single relation plans
- In each iteration greedily join to plans with the minimal cost
- Until a plan for the whole query has been generated

Other join enumeration techniques

- Randomized algorithms
 - randomly rewrite the join tree one rewrite at a time; use **hill-climbing** or **simulated annealing** strategy to find optimal plan.
- Genetic algorithms

Summary: (Join) Optimization

- Find “best” query execution plan based on a cost model (considering I/O cost, CPU cost,...); data statistics (histograms); dynamic programming, greedy join enumeration; physical plan properties (interesting orders).