

CS525: Advanced Database Organization

Notes 8: Concurrency Control

Gerald Balekaki

Department of Computer Science

Illinois Institute of Technology

gbalekaki@iit.edu

July 31st, 2023

Slides: adapted from a course taught by [Shun Yan Cheung](#), [Emory University](#)

Concurrency and Recovery

- DBMS should enable reestablish correctness of data in the presence of failures
 - System should restore a correct state after failure (recovery)
- DBMS should enable multiple clients to access the database concurrently
 - This can lead to problems with correctness of data because of interleaving of operations from different clients
 - System should ensure correctness (concurrency control)

Techniques to implement transactions

- Logging
 - Implements the **atomicity** property
 - Implements the **durability** property
- Synchronization (e.g.: locking)
 - Implements the **isolation** property
- The **consistency** property is assumed — otherwise, there is a bug in the transaction

Concurrent execution of transactions

- Concurrent execution of transactions can cause the database state to become inconsistent
- Example: Consider the following 2 transactions
 - T_1 : transfer \$1 from B to A
 - T_2 : transfer \$2 from B to A

T_1 :

```
READ(A)
A = A + 1
WRITE(A)
READ(B)
B = B - 1
WRITE(B)
```

T_2 :

```
READ(A)
A = A + 2
WRITE(A)
READ(B)
B = B - 2
WRITE(B)
```

Serial (non-concurrent) execution

Initial data value:

A = 10

B = 10

A + B = 20

T₁:

READ(A)

A = A + 1

WRITE(A)

READ(B)

B = B - 1

WRITE(B)

A = 11

B = 9

A + B = 20

T₂:

READ(A)

A = A + 2

WRITE(A)

READ(B)

B = B - 2

WRITE(B)

A = 13

B = 7

A + B = 20

Concurrent execution

```
Initial data value:
  A = 10
  B = 10
                                     A + B = 20
-----
T1:  READ(A)    (10)
      A = A + 1  (11)
      WRITE(A)   ==> A = 11
      READ(B)    (8)
      B = B - 1
      WRITE(B)   ==> B = 7
-----
T2:  READ(A)
      A = A + 2
      WRITE(A)   ==> A = 12
      READ(B)
      B = B - 2
      WRITE(B)   ==> B = 8
-----
Final data value:
  A = 11
  B = 7
                                     A + B ≠ 20
-----
```

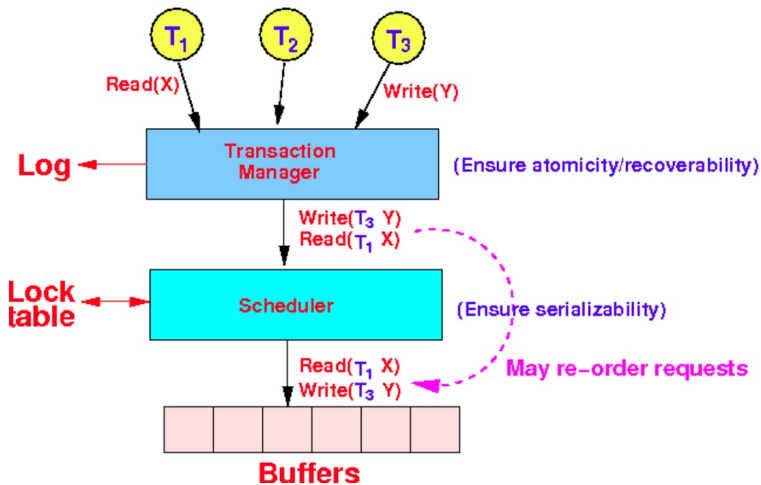
Observed fact from the above example

- The ordering of operations (actions) of different **transactions** can affect the **consistency** of the database
- To ensure **consistency** of the data:
 - We must impose a certain ordering on the operations from different **transaction**
 - Note: we still need to define what this “ordering” is, but we do know that we need some kind of ordering just by looking at the two examples

Scheduler and concurrency control

- **Scheduler**: the software component in a DBMS that is responsible for
 - regulating (scheduling) the executions of the actions (**read/write** operations) of the transactions
- **Concurrency control**: the procedure/algorithm to ensure/guarantee that:
 - a **concurrent** execution of the actions (operations) in (**multiple**) transactions will preserve the **consistency** of the database

Software organization of a DBMS



Software organization of a DBMS

- Transactions submit (read/write) requests for database elements to the Transaction Manager
- The transaction manager will write log records and then forward the (read/write) requests to the Scheduler
- The Scheduler can grant or delay the (read/write) requests based on the currently granted requests
 - If the read/write request is safe, the scheduler will grant the read/write request and the transaction can continue
 - If the read/write request is unsafe, the scheduler will delay the request and the transaction will wait
- A commonly used technique to decide on grant/delay is Locking
- Note:
 - In some concurrency control techniques (e.g., locking), the scheduler may need to abort transactions
 - This happens when transactions are dead locked

Topics that will be discussed

1. Define the correctness of a concurrent execution of transactions
 - This definition is called **serializability**
 - **Serializable schedule** (order of execution) will guarantee data consistency
2. **Serializability** turns out to be impractical to implement
 - We will look at a subset of correct schedules called **Conflict Serializability**
 - This notion of correctness can be implemented using **locks**
3. We will study how to implement **concurrent execution** of transactions are **conflict serializable** by using various kinds of locks
 - Exclusive locks
 - Shared/exclusive locks
 - Shared/update/exclusive locks
 - Shared/increment/exclusive locks

Serial schedules

- Transaction will transform a database from a consistent state to another consistent state if:
 - There are no system failures
 - The transaction is executed in isolation (no other transaction is active)
- Correctness of an execution of multiple transactions
 - Multiple transactions will transform a database from a consistent state to another consistent state if:
 - There are no system failures
 - Each transaction is executed in isolation (no other transaction is active)

- **Schedule:** a sequence of (important) operations performed by one or more transactions
- In concurrency control, the important operations performed by transactions are:
 - read operations (**READ**(X))
 - write operations (**WRITE**(X))

Schedule: Example

We have 2 DB elements:

A = 25

B = 25

(A = B)

We have 2 transactions:

T₁ : add 100 to database elements A and B

T₂ : double the value of database elements A and B

In code:

T ₁	T ₂
=====	=====
READ(A,t)	READ(A,s)
t = t + 100	s = 2×s
WRITE(A,t)	WRITE(A,s)
READ(B,t)	READ(B,s)
t = t + 100	s = 2×s
WRITE(B,t)	WRITE(B,s)

Example of a schedule:

```
READ1(A,t)
READ2(A,s)
t = t + 100
s = 2×s
WRITE1(A,t)
WRITE2(A,s)
READ1(B,t)
t = t + 100
WRITE1(B,t)
READ2(B,s)
s = 2×s
WRITE2(B,t)
```

Correctness consideration when executing T_1 and T_2

When T_1 or T_2 are executed in isolation:

T_1

```
READ(A,t)
t = t + 100
WRITE(A,t)
READ(B,t)
t = t + 100
WRITE(B,t)
```

T_2

```
READ(A,s)
s = 2 × s
WRITE(A,s)
READ(B,s)
s = 2 × s
WRITE(B,t)
```

T_1 and T_2 will preserve the following property:

$A = B$

Serial schedule

- **Serial schedule**: a schedule where the all operations of each transactions are executed in sequence
- i.e., operations from different transactions are not interspersed with one another
- Example:

```
Execute T1 then T2: // Subscript indicate the transaction
READ1(A,t)
t = t + 100
WRITE1(A,t)
READ1(B,t)
t = t + 100
WRITE1(B,t)

READ2(A,s)
s = 2 × s
WRITE2(A,s)
READ2(B,s)
s = 2 × s
WRITE2(B,t)
```


- Representation of a serial schedule
 - We can represent a serial schedule by a list of the transaction IDs.
 - T₁ T₂

Database consistency and serial schedules


- A **serial schedule** executes each **transaction** in isolation
- Therefore, every **serial schedule** will preserve the **consistency** (**correctness**) of the database
- Conclusion: We can use a **serial schedule** as the basis for the definition of correctness for **Concurrent execution of transactions**

Serializable schedule


- A schedule **S** is **serializable schedule** iff:
 - There is a **serial schedule S'** such that the effect of the execution of **S** and **S'** are identical for every **DB state**
- Starting in any **DB state**, executing the **schedules S/S'** will result in the same **DB state**



Example: serializable schedule




The interleaved execution (schedule) is equivalent to the following serial schedule:



The schedule is therefore serializable

Non-serializable schedules

- Consider the following schedule:



Non-serializable schedules

- The execution will result in an inconsistent database state:




- This result can never be obtained by a serial execution of T₁ and T₂
- Therefore, the schedule is not serializable

Determinant of serializable behavior


- To determine whether an interleaved execution of transaction is serializable, we must discard the semantics (operations performed) of the transaction
- The serializable of an interleaved execution of transaction is only determined by the interleaving of their `READ(...)` and `WRITE(...)` operations

Example 1

The schedule of T_1 and T_3 :




The schedule without semantics:




Example 2

The schedule of T_1 and T_2 :



The schedule without semantics:




The schedules without semantics (meaningful operations) in examples 1 and 2 are identical.

Determinant of serializable behavior

- The criteria to decide if a `schedule` is
 - `Serializable` or
 - `Not serializable`
- must be based solely on the ordering of the `READ(.)` and `WRITE(.)` operations in the `schedule`

Notation for a schedule

- $r_i(X)$: transaction T_i reads the DB element X
- $w_i(X)$: transaction T_i writes the DB element X
- Example: notation for a schedule



- Notation: $r_1(A) \ w_1(A) \ r_2(A) \ w_2(A) \ r_2(B) \ w_2(B) \ r_1(B) \ w_1(B)$

Conflict-serializability: a more practical type of serializability

- General definition of serializability:
 - S is **serializable** iff:
 - For every DB **state**, the effect of the execution of S is **identical/equivalent** to some **serial schedule**
- The general definition of **serializability** is difficult to apply (detect, verify).
- We have therefore define different kinds of **serializability** that are more practical (easier to detect):
 - **Conflict-serializability**: based on the concept of conflicting operations (This type of **serializability** can be easily enforced by **locks**)
 - **View-serializability**: based of the concept of view (what did the **transaction** “see” (read)) (This type of **serializability** is more general but harder to enforce)
- The most common form of **serializability** implemented in DMBS is **Conflict-serializability**

Conflicting operations

- 2 consecutive operations in a **schedule** S ($S = \dots \text{op1 op2} \dots$) is conflicting iff the behavior/result of at least one operation is changed when their order is interchanged

Interactions between the basic operations

- Recall: The basic operations used by transactions to operate on the database are `READ(.)` and `WRITE(.)`
- A `READ(.)` and `WRITE(.)` operation can operate:
 - On the same database elements
 - On different database elements

Non-conflicting and conflicting operations

- The interaction of 2 basic operations is **non-conflicting** if:
 - The ordering of the execution of the basic operation does not affect the outcome
 - Otherwise, the operations are **conflicting**
- For a **READ(.)** operation, the outcome is the return value of the read
- For a **WRITE(.)** operation, the outcome is the final value of the **database element** that is written

Read/write operations within the same transaction

- Read/write operations within a transaction will be executed in a certain order
- We assume that the ordering of read/write within one transaction is fixed and cannot be altered
- I.e., Operations from the same transaction can not be executed in a different order

Summary on conflicting operations

- 2 actions from the same transaction are conflicting
 - 2 actions from different transactions are conflicting if
 - They operate on the same DB element
 - At least one of the operation is **WRITE**
1. $r_i(X); w_j(X)$
 2. $w_i(X); r_j(X)$
 2. $w_i(X); w_j(X)$

Re-ordering non-conflicting operations

- If 2 actions `op1` and `op2` are **non-conflicting**, we can:
 - **Re-order** their execution in a **schedule** (and the result of the execution will remain unchanged)

Conflict-serializable schedules: conflict-serializability

- Conflict-equivalent schedules

- 2 schedules S_1 and S_2 are conflict equivalent iff

- Schedule S_1 can be transformed into schedule S_2 by a sequence of non-conflicting swaps of adjacent actions

- Example

- Consider the following 2 schedules

S_1 : $r_1(A)$ $w_1(A)$ $r_2(A)$ $w_2(A)$ $r_1(B)$ $w_1(B)$ $r_2(B)$ $w_2(B)$

S_2 : $r_1(A)$ $w_1(A)$ $r_2(A)$ $r_1(B)$ $w_2(A)$ $w_1(B)$ $r_2(B)$ $w_2(B)$

- The schedules are conflict equivalent because:

S_1 : $r_1(A)$ $w_1(A)$ $r_2(A)$ $w_2(A)$ $r_1(B)$ $w_1(B)$ $r_2(B)$ $w_2(B)$

S_2 : $r_1(A)$ $w_1(A)$ $r_2(A)$ $r_1(B)$ $w_2(A)$ $w_1(B)$ $r_2(B)$ $w_2(B)$

Conflict-serializable schedules: conflict-serializability

- Conflict-serializable schedules

- A schedule S_1 is conflict serializable iff

- Schedule S_1 can be transformed into a serial schedule by a sequence of non-conflicting swaps of adjacent actions

- Example

- Consider the following schedule

S_1 : $r_1(A)$ $w_1(A)$ $r_2(A)$ $w_2(A)$ $r_1(B)$ $w_1(B)$ $r_2(B)$ $w_2(B)$

- We perform the following sequence of non-conflicting swaps and obtain a serial schedule

S_1 : $r_1(A)$ $w_1(A)$ $r_2(A)$ $w_2(A)$ $r_1(B)$ $w_1(B)$ $r_2(B)$ $w_2(B)$

S_1 : $r_1(A)$ $w_1(A)$ $r_2(A)$ $r_1(B)$ $w_2(A)$ $w_1(B)$ $r_2(B)$ $w_2(B)$

S_1 : $r_1(A)$ $w_1(A)$ $r_1(B)$ $r_2(A)$ $w_2(A)$ $w_1(B)$ $r_2(B)$ $w_2(B)$

S_1 : $r_1(A)$ $w_1(A)$ $r_1(B)$ $r_2(A)$ $w_1(B)$ $w_2(A)$ $r_2(B)$ $w_2(B)$

S_1 : $r_1(A)$ $w_1(A)$ $r_1(B)$ $w_1(B)$ $r_2(A)$ $w_2(A)$ $r_2(B)$ $w_2(B)$

S_1 : T_1 T_2

- \therefore The schedule S_1 is a conflict-serializable schedule

Conflict-serializable schedule are always serializable

- Swapping two non-conflicting operations in a schedule will not change the effect of the execution of the operations in a schedule
- Therefore, the effect of the execution of a conflict-serializable schedule is equivalent to the execution of a serial schedule of the transactions
- So by definition of serializability, a conflict-serializable schedule is a serializable schedule

Schedules can be serializable but not conflict-serializable

- Example of a schedule that is **serializable** but **non-conflict serializable**
 - Consider the following 3 transactions
 - $T_1: \quad w_1(Y) \quad w_1(X)$
 - $T_2: \quad w_2(Y) \quad w_2(X)$
 - $T_3: \quad w_3(X)$
 - Consider the following 2 schedules with the 3 transactions: (Schedule S_2 is a **serial schedule**)
 - $S_1: \quad w_1(Y) \quad w_2(Y) \quad w_2(X) \quad w_1(X) \quad w_3(X)$
 - $S_2: \quad w_1(Y) \quad w_1(X) \quad w_2(Y) \quad w_2(X) \quad w_3(X)$
 - The effect (state of the database) of schedules S_1 and S_2 are identical
 - The database state: the updated made by the last write operation
 - By definition: **Schedule S_1** is equivalent to a serial schedule
 - However, we cannot transform the **schedule S_1** into a serial schedule by swapping adjacent **non-conflicting actions**
 - \therefore The **schedule S_1** is not a **conflict-serializable schedule**

Precedence graph test for conflict-serializability

- The reason why do we use **conflict-serializability**, there is a simple procedure the test/check whether a **schedule S** is **conflict-serializable** or not **conflict-serializable**
- The precedence relationship
- Observation
 - If we find 2 **conflicting** actions (anywhere in a schedule)



- then T_i must precede T_j in a **serial schedule** (because $op_j(X)$ can never be swap with $op_i(X)$)
- Therefore, the ordering in a **serial schedule** must be as follows




Precedence graph test for conflict-serializability

- This observation has determined a **precedence relationship** between transactions T_i and T_j .
 - Specifically, T_i must precede T_j (in a **serial schedule**)
- Representing (all) the precedence relationships in a **schedule S**
 - We use a graph to represent all **precedence relationships** found in a **schedule S**
 - For every pair of conflicting actions in a **schedule S** we add the edge $T_i \rightarrow T_j$ to the graph

Example

- S_1 : $r_2(A) \ r_1(B) \ w_2(A) \ r_3(A) \ w_1(B) \ w_3(A) \ r_2(B) \ w_2(B)$
- Graph with all precedence relationships
 - Specifically, T_i must precede T_j (in a serial schedule)
- Representing (all) the precedence relationships in a schedule S




A simple test for conflict-serializable schedules

- A schedule S is conflict-serializable iff the corresponding precedence graph of schedule S does not contain any cycles
- Example:

S_1 : $r_2(A)$ $r_1(B)$ $w_2(A)$ $r_3(A)$ $w_1(B)$ $w_3(A)$ $r_2(B)$ $w_2(B)$

- Graph with all precedence relationships




- The serial schedule is: T_1 T_2 T_3
- Proof: swap the non-conflicting operations

A simple test for conflict-serializable schedules:

Example 2

- S_1 : $r_2(A) \ r_1(B) \ w_2(A) \ r_2(B) \ r_3(A) \ w_1(B) \ w_3(A) \ w_2(B)$
- Graph with all precedence relationships



- This schedule is not conflict-serializable because there is a cycle
- Proof
 - To make a serial schedule with the order $T_1 \dots T_2$ we must swap this pair of conflicting action:
 S_1 : $r_2(A) \ r_1(B) \ w_2(A) \ r_2(B) \ r_3(A) \ w_1(B) \ w_3(A) \ w_2(B)$
 - To make a serial schedule with the order $T_2 \dots T_1$ we must swap this pair of conflicting action:
 S_1 : $r_2(A) \ r_1(B) \ w_2(A) \ r_2(B) \ r_3(A) \ w_1(B) \ w_3(A) \ w_2(B)$
 - Therefore: there is no way to swap non-conflicting actions to convert the schedule into a serial schedule

- Overview: techniques to enforce `conflict-serializable` behavior

Techniques to enforce conflict-serializable behavior

- Locks
 - Exclusive locks
 - Shared/Exclusive locks
 - Shared/Update/Exclusive locks
 - Shared/Increment/Exclusive locks
- Locking will also require 2-phase locking to be sufficient (later)
- Timestamps
 - Single-value timestamp
 - Multi-value timestamp
- Validation
 - Verification

Lock

- Lock: an object that have:
 - A state (locked, unlocked, etc)
 - A number of lock holders
 - A queue of pending (locking/unlocking) operations
- Some lock operation will be **granted** (succeed) when the lock is in a certain state
 - If a lock operation is **granted**, the execution will continue
- Some lock operation be **denied** (fail) when the lock is in a certain state
 - If a lock operation is **denied**, the execution will wait until the state of the lock changes to a compatible state
 - When the state of the lock becomes compatible, the lock operation will be **granted** and the execution will be continued

Transactions, conflict-serializability and locks

- Uncontrolled access to DB elements will result in non-serializable behavior
- Locks can alter the order of the execution of operations
 - By denying a request for a lock, we can alter the order of the execution of certain operations
- Goal
 - Define grant/deny rules so that the modified order of operations is conflict-serializable

Exclusive locks

- Lock states
 - unlocked
 - (exclusive) locked
- Lock operations
 - `lock`
 - `lock requests` is granted if the `lock state` is `unlocked`
 - The state of the lock becomes `locked`
 - `lock requests` is denied if the `lock state` is `locked`
 - `unlock`
 - The state of the lock will become `unlocked`
 - Only applicable if the transaction currently hold the lock
- Notations
 - $l_i(X)$: transaction `i` requests a lock on DB element `X`
 - $u_i(X)$: transaction `i` requests to unlock the DB element `X`

Locking rules for exclusive locks

- Rules


1. When a transaction T want to perform an action (read/write) on a DB element X , transaction T must first obtain the lock on DB element X
2. The transaction T must release all locks before it terminates

- $T_i: \dots \quad l_i(X) \quad r_i(X) \quad \dots \quad u_i(X) \quad \dots(\text{read})$
- $T_i: \dots \quad l_i(X) \quad w_i(X) \quad \dots \quad u_i(X) \quad \dots(\text{write})$

- Note: transactions must hold the lock to perform an action. So this is not allowed

- $\dots \quad l_i(X) \quad \dots \quad u_i(X) \quad w_i(X) \quad \dots(\text{write})$

The lock scheduler



- **Lock scheduler:** a scheduler that uses locks to make re-ordering decisions in the execution of the actions of the **transactions**
 - The **scheduler** will grant the **lock request** for an operation if the request cannot result in an **inconsistent state**
 - The **scheduler** will deny the **lock request** for an operation if the request can result in an **inconsistent state**

Lock table: an intro

- ~~The exact implementation of a lock table will be discussed later~~
- For now, imagine that a **lock table** is a set of tuples of the following format:
 - (DB Element, Transaction holding lock)
- Example:
 - (X, T_1)
 - (Y, T_2)
- Note: Due to the locking rule of exclusive locks, only one transaction can hold a lock

Example of locking

- Transactions

T_1 :

$r_1(A)$

$A = A + 1$

$w_1(A)$

$r_1(B)$

$B = B - 1$

$w_1(B)$

T_2 :

$r_2(A)$

$A = 2 \times A$

$w_2(A)$

$r_2(B)$

$B = 2 \times B$

$w_2(B)$

- Transactions augmented with lock operations

T_1 :

$l_1(A)$

$r_1(A)$

$A = A + 1$

$w_1(A)$

$u_1(A)$

$l_1(B)$

$r_1(B)$

$B = B - 1$

$w_1(B)$

$u_1(B)$

T_2 :

$l_2(A)$

$r_2(A)$

$A = 2 \times A$

$w_2(A)$

$u_2(A)$

$l_2(B)$


$r_2(B)$

$B = 2 \times B$

$w_2(B)$

$u_2(B)$

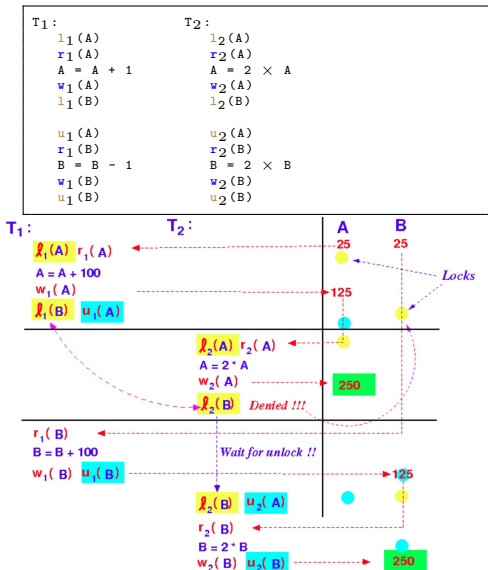
Example of a schedule with locks



- Notice that
 - The transactions obeys the locking rules
 - ... $l_i(X) r_i(X) \dots w_i(X) u_i(X) \dots$
 - However, the resulting schedule is **non-serializable**.
 - (Because serial schedules of transactions T₁ and T₂ will preserve the property A = B)
- Conclusion: The locking rules is insufficient to ensure (conflict)-serializability

Example 2 of locking (with a correct outcome)

Consider a slightly changed order in the transactions: We re-order these 2 unlock/lock operations in the previous example:



Example 2 of locking (with a correct outcome)

- Conclusion
 - The execution results in a **consistent state** $A = B$
 - So **locks** can force a **schedule** to change the order of execution of the actions so that the resulting execution of the **schedule** is **conflict-serializable**
- Question: What is the secret ingredient that we need to add to make **locking** a sufficient technique to ensure **conflict-serializability**


Two Phase Locking (2PL)

- **Two Phase Locking:** ordering of locking operations in a transaction where all the lock operations **precedes** all the unlock operations


```
Transaction:
    lock(A)
    ...
    lock(B)
    ...
    ...
    lock(X)
    ...
    ...
    unlock(B)
    ...
    unlock(X)
    ...
    unlock(A)
    ...
```

- The 2 phases
 - Phase 1: Acquire locks
 - Phase 2: Release locks

Two Phase Locking (2PL): Example




- Notice: All lock operations **precedes** all unlock operations
- We saw in the previous example (slides 53) that execution using 2-phase locking results in a consistent database state



Two Phase Locking (2PL)

- Recall: Transactions using non 2-phase locking can result in a inconsistent DB state
- As in previous example (slides 51, 52), the transactions are not two-phase since the unlock operation $u_i(A)$ happens before an lock operation $l_i(B)$
- Results in execution results in inconsistent database state



- Proof: 2PL is sufficient condition for conflict-serializability (omitted)

Deadlock


- **Deadlock:** a state of execution where two or more transactions are unable to proceed because the transactions involved are waiting of each other's locks
- Example

Consider the following 2 transactions:

T ₁ :	T ₂ :
$l_1(A)$	$l_2(B)$
$r_1(A)$	$r_2(B)$
$A = A + 1$	$B = 2 \times B$
$w_1(A)$	$w_2(B)$
$l_1(B)$	$l_2(A)$
$u_1(A)$	$u_2(B)$
$r_1(B)$	$r_2(A)$
$B = B - 1$	$A = 2 \times A$
$w_1(B)$	$w_2(A)$
$u_1(B)$	$u_2(A)$

Deadlock: Example

- A possible interleaved execution (that obeys the locking rules) is




2PL and deadlocks

- 2PL can not prevent **deadlocks**
 - In the previous example (slides 58,59), both **transactions** are 2PL.
 - However, their execution resulted in **deadlock**
- Deadlock detection
 - using time out
 - using wait-for graph
- Deadlock prevention
 - using ordered DB elements
 - using wait-die timestamp scheme
 - using wound-wait timestamp scheme

Shared/exclusive locks


- Consider the following 2 transactions



- Notice that: T₁ and T₂ only reads the DB elements A and B

Shared/exclusive locks

- A possible schedule of T_1 and T_2



- T_2 could proceed when T_1 holds the lock for DB element A because the value of the DB element A will not be changed by T_1 . So there was no technical reason to block the T_2
- The exclusive lock is too restrictive in some situations

Shared/exclusive locks: locking method

- Uses 2 kinds of locks:
 - Shared lock
 - Exclusive lock
- DB element X can have one of the following locks
 - 1 exclusive lock
 - any number of shared locks
- Transactions using shared/exclusive locking
 - In order to perform a read operation on a DB element X, the transaction must be holding a shared lock on the DB element X
 - In order to perform a write operation on a DB element X, the transaction must be holding an exclusive lock on the DB element X
 - A shared/exclusive lock operation must be followed by an unlock operation

Shared/exclusive locks: Notation

- $sl_i(X)$: transaction T_i requests a shared lock on the DB element X
- $xl_i(X)$: transaction T_i requests an exclusive lock on the DB element X
- $u_i(X)$: transaction T_i unlocks on the DB element X
 - $u_i(X)$ can only be applied if transaction T_i currently holds a (shared or exclusive) lock on the DB element X


Requirements to ensure conflict-serializability

- The following 3 types of requirements can guarantee that current execution transactions will be conflict-serializable
 - Transaction consistency (semantics)
 - In order to perform a read operation on a DB element X, the transaction must be holding a shared lock on the DB element X
 - In order to perform a write operation on a DB element X, the transaction must be holding an exclusive lock on DB element X
 - A shared/exclusive lock operation must be followed by an unlock operation
 - 2 Phase Locking: All the unlock operations must follow all the lock operations within one transaction
 - Legality of the lock schedule
 - A DB element X can be in one of the following states
 - unlocked
 - exclusively locked by one transaction (with no shared locks)
 - shared locked by one or more transactions

- A transaction cannot hold on the same DB element
 - a shared lock and
 - an exclusive lock
- Therefore, a transaction that read the DB element X and then write the DB element X must request an exclusive lock for the DB element X

Lock compatibility matrix

- If we use several lock modes, then the scheduler needs a policy about when it can grant a lock request, given the other locks that may already be held on the same database element
- There is a clearer way to represent the **compatibility** of locks by using a **compatibility matrix**
- Structure of a compatibility matrix



- State of a lock: the type of lock that has been granted
- Lock request type: the type of lock being requested
- The compatibility decision is
 - Yes, if the requested lock type can be granted
 - No, if the requested lock type can not be granted

Lock compatibility matrix for shared/exclusive locks

- Full compatibility matrix for shared/exclusive locks:

<i>State of the lock</i>	<i>Lock request type</i> →	
	<i>Shared</i>	<i>Exclusive</i>
<i>Unlock</i>	Yes	Yes
<i>Shared</i>	Yes	No
<i>Exclusive</i>	No	No


← *Trivial*

- Abbreviated compatibility matrix for shared/exclusive locks:


<i>State of the lock</i>	<i>Lock request type</i> →	
	<i>Shared</i>	<i>Exclusive</i>
<i>Shared</i>	Yes	No
<i>Exclusive</i>	No	No

Upgrading locks

- Technically it's very easy to implement upgradable locks
 - A transaction **T** that wants to read a DB element first acquires a **shared lock**
 - We set the state of the lock to **SHARED**
 - When transaction **T** wants to write a DB element later, it can update the **shared lock** to an **exclusive lock**
 - We simply set the state of the lock to **EXCLUSIVE**
- Example



Example of interleaved execution (possible schedule)




The danger of upgrading locks

- lock upgrade can result in deadlock with just one DB element

Consider the following 2 transactions:

T ₁ :	T ₂ :
sl ₁ (A)	sl ₂ (A)
r ₁ (A)	r ₂ (A)
xl ₁ (A)	xl ₂ (A)
w ₁ (A)	w ₂ (A)
u ₁ (A)	u ₂ (A)

- Possible interleaved execution:



Increment lock

- The increment operation and the increment lock
 - The increment operation: **INC(A, c)**: performs the following steps atomically

```
1 READ(A, t) //t is a local variable in the tranx
2 t = t + c
3 WRITE(A, t)
```


Note: c can be negative

- Many applications make use of the increment operation
 - banking: transfer sum
$$A = A + c$$
$$B = B + (-c)$$
 - booking
$$\text{\#seats} = \text{\#seats} + (-1)$$

Increment lock

- Types of operations on database elements
 - `READ(.)`
 - `WRITE(.)`
 - `INC(., c)`
- For each type of operation, we use a specific type of lock
 - **shared lock**
 - grants a transaction the privilege to perform a **read operation** on the **DB element**
 - **exclusive lock**
 - grants a transaction the privilege to perform a **write operation** (also **read**) on the **DB element**
 - **increment lock**
 - grants a transaction the privilege to perform an **increment operation** on the **DB element**
- $INC_i(A, c_1)$ and $INC_j(A, c_2)$ do not conflict

Compatibility matrix for shared/increment/exclusive locking

<i>State of the lock</i>	<i>Lock request type</i> 		
	<i>Shared</i>	<i>Exclusive</i>	<i>Increment</i>
<i>Shared</i>	Yes	No	No
<i>Exclusive</i>	No	No	No
<i>Increment</i>	No	No	Yes